

EE 241B HW2 Writeup

Vighnesh Iyer

Contents

1	Extracting and Simulating a Synthesized Design	1
1.1	Delay of a Path	1
1.2	ICC Critical Path + Special Path	3
1.3	Power Estimate Accuracy Analysis	3
1.4	Voltage Scaling Power Estimates	5
2	FF and Latch Based Timing	6
2.1	FF Timing	6
2.2	Latch Timing Phases	6
2.3	Latch Timing Clock Period	6
3	Variability and Timing Simulations	7
3.1	Delay Variation Across Process Corners	8
3.2	Monte Carlo Simulation for Transistor Mismatch	8
3.3	Delay Variation with Temperature	9
3.4	Impact of Variability on Timing	10

1 Extracting and Simulating a Synthesized Design

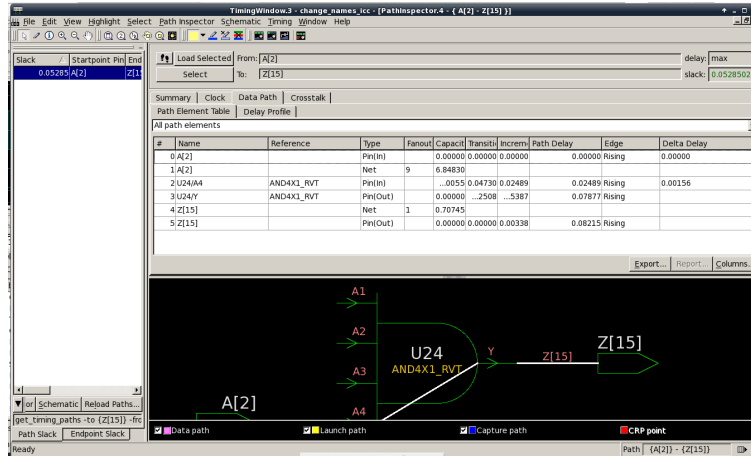
We start with an already placed and routed 4-bit decoder and we go through the flow of extracting parasitics, running LVS, and running SPICE simulations of the extracted schematic.

1.1 Delay of a Path

For the path from A[2] to Z[15] rise, what is the delay measured by IC Compiler, SPICE simulation without extracted parasitics, and SPICE simulation with parasitics?

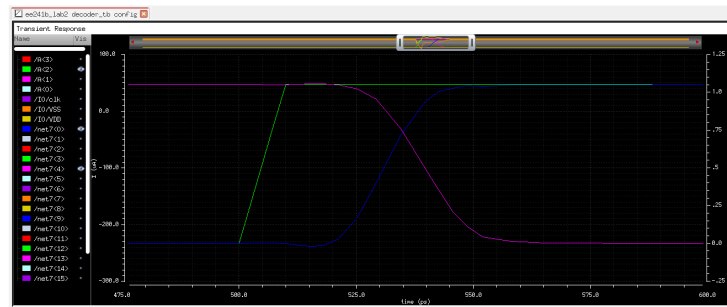
- IC Compiler: 0.08215 ns = 82.15 ps

This number was derived from the placed and routed design loaded into ICC and the timing path inspector was used to extract the relevant path and display its delay.



- SPICE without parasitics: 26.297 ps


This number was derived by importing the layout and Verilog netlist into Virtuoso, performing LVS to check that the schematic and layout were representative of the same circuit, and then using the schematic to perform a transistor level SPICE simulation of rising A[2] to rising Z[15].



In the figure you can see the rise of A[2] (green) followed by the rise of Z[15] (blue) and the fall of Z[11] (pink). The delay in this part is significantly lower than the delay as reported by ICC; it could have to do with parasitics, which will be simulated in the next part.

- SPICE with parasitics: 46.91 ps

This number was derived using the same testbench, but using the 'starcc' parasitic extracted view of the decoder schematic rather than the raw schematic with ideal circuitry. We can see that the delay increased by almost 2x due to parasitics.

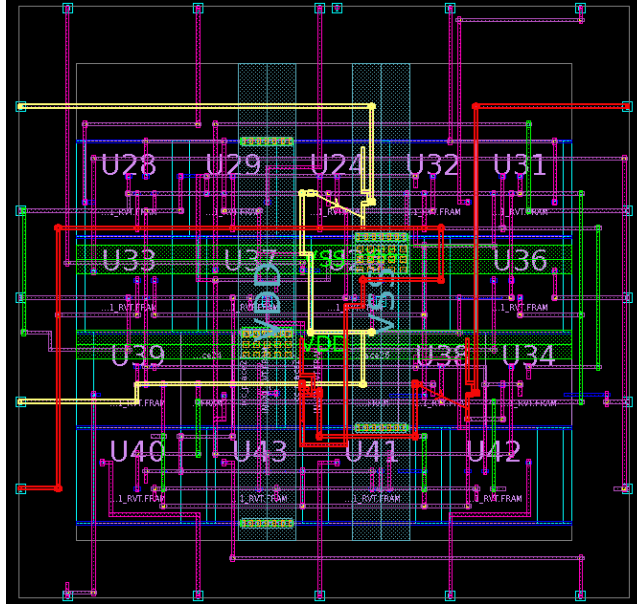


This figure has the same color scheme as the previous one. The path delay is still roughly half of what ICC reported. This probably has to do with some inherent pessimism in the tool.

1.2 ICC Critical Path + Special Path

We display the critical path in ICC and also highlight the special path noted above from the rise of A[2] to the rise of Z[15].

The critical path is from rising A[3] to falling Z[5], and it has a negative slack of -0.00263 ns. It is highlighted in red. The special path above is highlighted in yellow.



1.3 Power Estimate Accuracy Analysis

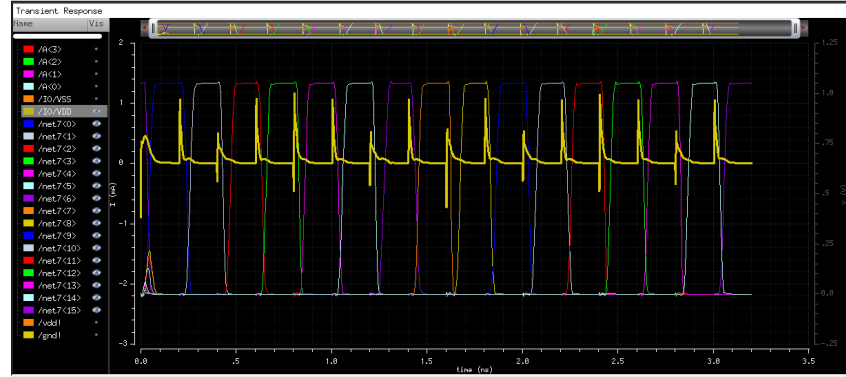
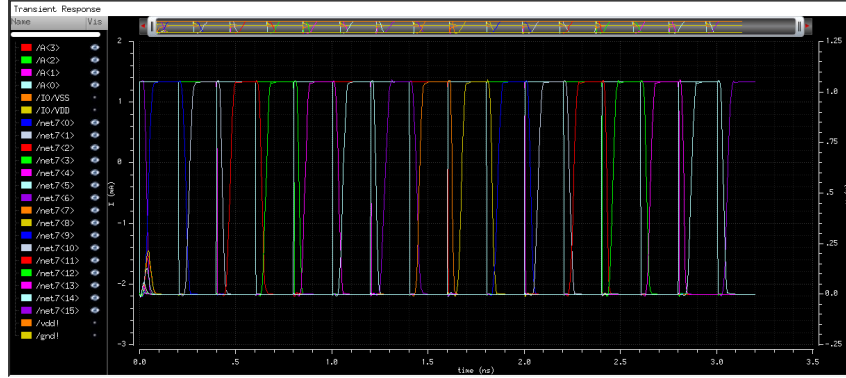
We run the functional testbench that counts from A=0000 to A=1111 and report the power measured by Primetime, the SPICE simulation, and the mixed-signal simulation.

- Primetime: 3.97e-5 (switching), 4.35e-5 (int), 1.90e-6 (leakage), 8.51e-5 (total) = 85 uW

We run a post-PAR simulation using VCS and extract the switching activity SAIF file from the output of the testbench. This file and the Verilog netlist is fed into Primetime which estimates the power usage for the testbench.

- SPICE: 73.46 uW

We re-create the Verilog testbench in SPICE by using pulsed voltage sources for the `A<3:0>` input to the decoder and we measure the average power consumption by integrating the current drawn from the VDD supply during the transient SPICE simulation.




The first plot shows the switching activity of both the inputs and the outputs of the decoder. The second plot shows the current drawn from the VDD input (bold yellow line) of the decoder and each output bit of the decoder going high as its respective code is applied to the decoder's input. The power consumed is measured by using the **average** function on the transient current on the decoder's VDD port and multiplying by the supply voltage of 1.05 V.

The power consumed as reported by SPICE with parasitics is a little lower than what was reported by PrimeTime.

- Mixed-Signal: 75.79 uW

We then set up a mixed-signal testbench by exporting the SPICE netlist of the extracted decoder. We create a wrapper for the decoder in Verilog and instantiate it in a testbench. The testbench is used to drive the inputs and check the outputs of the decoder, while the

actual function of the decoder is simulated at the transistor level. This combines the best parts of the Primetime and SPICE flows by letting us write stimulus at a high-level (in Verilog), but still retain the precision of a transistor level simulation using SPICE.



This snapshot from DVE shows the digital stimulus applied to the decoder and the SPICE simulated decoder's analog outputs.

1.4 Voltage Scaling Power Estimates

We run the mixed-signal simulation at 1.05V, 0.8V, 0.6V, and 0.4V and measure the average power for each voltage. We then convert average power to energy/op in terms of J/op and uW/Mhz. These results are compared to the theoretically predicted active energy savings for voltage scaling based on the 1.05V result.

The (fake) clock period will have to increase for lower supply voltages to compensate for increased delay.

We theoretically predict that the dynamic power scales according to this equation:

$$P_{dyn} = C_{eff} \cdot V_{DD}^2 \cdot f_{clk} \quad (1)$$


with C_{eff} being fixed for each supply voltage.

Here are my summarized results that were produced by varying the supply voltage to the decoder and measuring average power in the same manner. The number of operations is $2^4 = 16$ and $P_{scale,pred}$ is the predicted power scaling, P_{pred} is the predicted power consumption, and P_{sim} is the mixed-signal simulated actual power consumed. t_{sim} is used to measure the time of the testbench. The clock period was changed manually until signal settling was sufficiently fast.

V_{DD}	T_{clk}	$P_{scale,pred}$	P_{pred}	P_{sim}	t_{sim}	J/op	uW/Mhz
1.05 V	0.15 ns	1.0	75.79 uW	75.79 uW	2.4 ns	11.37 pJ	1.14e-8
0.8 V	0.17 ns	0.5122	38.82 uW	50.27 uW	2.72 ns	8.546 pJ	8.55e-9
0.6 V	0.5 ns	0.098	7.4274 uW	13.56 uW	8 ns	6.78 pJ	6.78e-8
0.4 V	decoder doesn't work (5ns)	0.0044	0.333 uW	10 uW	80ns	50 pJ	5e-8

2 FF and Latch Based Timing

We consider this logic path.



2.1 FF Timing

If registers R1 and R2 are positive edge triggered ($t_{setup} = t_{hold} = t_{clk-q} = 100$ ps), what is the minimum clock period that this system can operate at? Assume zero skew.

$$T_{clk} > t_{clk-q} + t_{logic,max} + t_{hold}$$

$$T_{clk} > 100 + 800 + 100 = 1000 \text{ ps}$$

2.2 Latch Timing Phases

If R1 and R2 are level sensitive, label the clock phases so that the system operates correctly.

If we are using fully out of phase clocks with no overlap, then R1 should be clocked by the general system clock, and R2 should be clocked with an inverted (180 ° out of phase) system clock.

2.3 Latch Timing Clock Period

We find the shortest clock cycle in a latch-based system with $t_{setup} = t_{hold} = t_{clk-q} = t_{d-q} = 100$ ps. The clock has a fixed 50% duty cycle and only 2 phases of the clock can be used for clocking. There is zero skew.

This rule must hold to prevent late latch arrivals:

$$T_{cycle} + PW > t_{clk-q} + t_{logic,max} + t_{setup} + t_{skew} \quad (2)$$

$$T_{cycle} + PW > t_{d-q} + t_{logic,min} + t_{borrowed} \quad (3)$$

We also enforce this rule to prevent premature latch arrivals (data races):

$$t_{clk-q} + t_{logic,min} > PW + t_{hold} + t_{skew} \quad (4)$$

PW is the pulse width and we set it to half the T_{cycle} for this problem.

Plugging in our values, we get:

$$1.5 \cdot T_{cycle} > 100 + 800 + 100 + 0$$

$$1.5 \cdot T_{cycle} > 100 + 50 + t_{borrowed}$$

$$100 + 50 > 0.5T_{cycle} + 100$$

$$1.5 \cdot T_{cycle} > 100 + 800 + 100 + 0$$

$$1.5 \cdot T_{cycle} > 100 + 50 + t_{borrowed}$$

$$100 + 50 > 0.5T_{cycle} + 100$$

3 Variability and Timing Simulations

We begin by recording the nominal (TT) FO4 low to high and high to low delays of an inverter. We size the inverters for minimum average delay.

Using SPICE and the 32nm LP model with a 1.05V supply, we find the optimal width of the PMOS transistor that minimizes the propagation delay ($t_p = (t_{pHL} + t_{pLH})/2$) for a CMOS inverter. We fix the NMOS transistor width at 100nm and the lengths of both transistors are fixed at the process' 32nm minimum.

We use the F04 (fanout of 4) method of inverter delay characterization. I designed a simulation based on the paper "The Fanout-of-4 Inverter Delay Metric" by Harris, Horowitz, et al. The circuit setup from the paper is shown below.


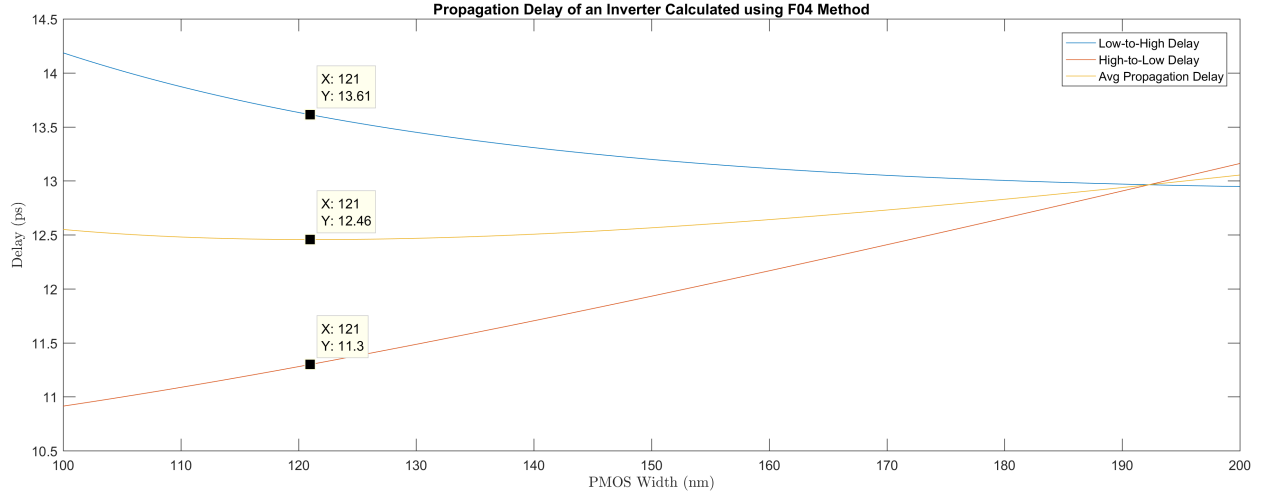


Figure 1: Fanout-of-4 inverter delays

I swept the width of the PMOS from 100n to 200n while performing transient simulations. I measured the propagation delays (high-to-low and low-to-high) using '.measure' statements in SPICE. Propagation delay was measured from 50% of the input to 50% of the output. The results are summarized below:



We find that the average propagation delay t_p is minimized when the PMOS width is 121nm and the minimal delay is 12.4560 ps. The PMOS width that results in an equal low-to-high and high-to-low delay is 192.3nm and the equal delay is 12.96 ps. In conclusion,:

$$W_{pmos,optimal} = 121 \text{ nm for } W_{nmos} = 100 \text{ nm}$$

$$t_{p,optimal} = 12.4560 \text{ ps}$$


3.1 Delay Variation Across Process Corners

We use the same F04 testbench and the optimally sized inverter, but now we vary process corners (SS, FF, SF, and FS) and log how the delay L-H and H-L is affected by the process corner.

Corner	Delay L-H (ps)	Delay H-L (ps)
TT	13.613	11.299
FF	9.422	7.877
SS	18.316	15.645
SF	13.926	14.514
FS	15.418	6.442

3.2 Monte Carlo Simulation for Transistor Mismatch


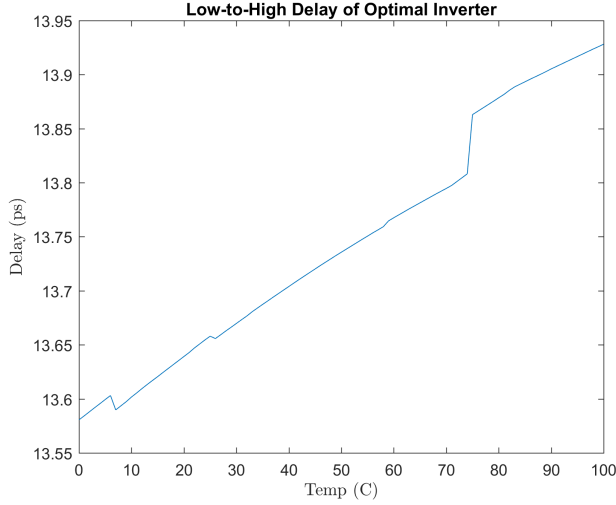
The transistor models use the Pelgrom mismatch model with $A_{Vt} = 2 \text{ mV}\mu\cdot\text{m}$. We run a 300-point Monte Carlo simulation and find the standard deviation of the delay, as well as the minimum and maximum delays observed.



The standard deviation of the L-H delay is 0.852 ps and H-L delay is 0.871. The range of L-H delay is (11.0543 ps, 16.7512 ps) and H-L delay is (9.5815 ps, 14.2203 ps).

3.3 Delay Variation with Temperature

We want to find how the delay (L-H and H-L) of an inverter changes with operation at 0° C to 100° C.



The collected data shows that temperature plays a much smaller role than process variation in the delay of an optimal inverter.

The L-H delay goes from 13.58 ps to 13.93 ps and the H-L delay goes from 11.28 ps to 11.71 ps.

3.4 Impact of Variability on Timing

If the critical path consisted of 7 F04 inverter stages, how would we set the nominal clock period? How much slower would the clock have to be set under worst-case global conditions and 3 sigma threshold variation?

We can use the nominal average delay of a F04 inverter stage as 12.46 ps. 7 of these stages gives us a total combinational delay of 87.22 ps. Thus the nominal clock period should be $87.22 + t_{clk-q} + t_{setup}$. Estimating $t_{clk-q} = t_{setup} = 20ps$, the nominal clock period is: 127.22 ps.

Under worst-case global conditions, we have to assume the slowest L-H propagation delay of 16.75 ps. Then we get a clock period of 157.25 ps.

Under 3 sigma threshold variation, we take the mean of the average propagation delay which is 12.59 ps and add 3 times the overall SD which is $3 \cdot 0.6036$ ps. Thus the clock period becomes: 140.81 ps.